

Principled Parsing for Indentation-Sensitive Languages

Revisiting Landin’s Offside Rule

Michael D. Adams

Portland State University
<http://michaeldadams.org/>

Abstract

Several popular languages, such as Haskell, Python, and F#, use the indentation and layout of code as part of their syntax. Because context-free grammars cannot express the rules of indentation, parsers for these languages currently use *ad hoc* techniques to handle layout. These techniques tend to be low-level and operational in nature and forgo the advantages of more declarative specifications like context-free grammars. For example, they are often coded by hand instead of being generated by a parser generator.

This paper presents a simple extension to context-free grammars that can express these layout rules, and derives GLR and LR(*k*) algorithms for parsing these grammars. These grammars are easy to write and can be parsed efficiently. Examples for several languages are presented, as are benchmarks showing the practical efficiency of these algorithms.

Categories and Subject Descriptors D.3.1 [Programming Languages]: Formal Definitions and Theory—Syntax; D.3.4 [Programming Languages]: Processors—Parsing; F.4.2 [Mathematical Logic and Formal Languages]: Grammars and Other Rewriting Systems—Parsing

General Terms Algorithms, Languages

Keywords Parsing, Indentation, Offside rule

1. Introduction

Languages such as Haskell [Marlow (ed.) 2010] and Python [Python] use the indentation of code to delimit various grammatical forms. In Haskell, the contents of a `let`, `where`, `do`, or `case` expression can be indented relative to the surrounding code instead of being explicitly delimited by curly braces. In Python, the body of a multi-line function or compound statement must be indented relative to the surrounding code; there is no alternative, explicitly-delimited syntax. For example, in Haskell one may write:

```
mapAccumR f = loop
  where loop acc (x:xs) = (acc'', x' : xs')
        where (acc'', x') = f acc' x
              (acc', xs') = loop acc xs
  loop acc [] = (acc, [])
```

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. To copy otherwise, to republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee.

POPL '13, January 23–25, 2013, Rome, Italy.

Copyright © 2013 ACM 978-1-4503-1832-7/13/01...\$15.00

The indentation of the bindings after each `where` keyword determines the parse structure of this code. For example, the indentation of the last line determines that it is part of the bindings introduced by the first `where` instead of the second `where`.

Likewise, in Python one may write:

```
def factorial(x):
    result = 1
    for i in range(1, x + 1):
        result = result * i
    return result
print factorial(5)
```

Here the indentation determines that the `for` loop ends before the `return` and the `factorial` function ends after the `return`.

While Haskell and Python are well known for being indentation-sensitive languages, quite a few other languages also use indentation. Landin’s ISWIM [Landin 1966] introduced the concept of the offside rule for indentation, which requires that all tokens in an expression be indented at least as far as the first token of the expression. Variations on this rule are used by Haskell, Miranda [Turner 1989], *occam* [INMOS Limited 1984], *Orwell* [Wadler 1985], Curry [Hanus (ed.) 2006], and *Habit* [HASP Project 2010]. F# is indentation sensitive when its lightweight syntax is enabled [Syme et al. 2010, §15.1]. The block styles in the *YAML* [Ben-Kiki et al. 2009] data serialization language are indentation sensitive, as are many forms in the *Markdown* [Gruber] and *reStructuredText* [Goodger 2012] markup languages. Even *Scheme* has an indentation-sensitive syntax in the form of *SRFI-49* [Möller 2005], though it is not often used.

Whitespace sensitivity may be controversial, but regardless of whether it is a good idea from a language design perspective, it is important that the grammars of layout-sensitive languages be precisely specified. Unfortunately, many language specifications are informal in their description of layout or use formalisms that are not amenable to practical implementation. The task of parsing layout is thus often left to *ad hoc*, handwritten code.

The lack of a standard formalism for expressing these layout rules and of parser generators for such a formalism increases the complexity of writing parsers for these languages. Often, practical parsers for these languages have significant structural differences from the language specification. For example, the layout rule for Haskell is specified in terms of an extra pass between the lexer and the parser that inserts explicit delimiters. This extra pass uses information about whether the parsing that occurs later in the pipeline will succeed or fail on particular inputs. Due to the resulting cyclic dependency, Haskell implementations do not actually structure their parsers this way. As a result, the structural differences between the implementation and the specification make it difficult to determine if one accurately reflects the other.

This paper aims to resolve this situation by proposing a grammar formalism for expressing layout rules. This formalism is both theoretically sound and practical to implement. Indentation-sensitive grammars are easy and convenient to write, and fast and efficient parsers can be implemented for them. The primary contributions of this paper are:

- a grammar formalism for expressing indentation-sensitive languages, which is informally described in Section 2 and formally defined later in Section 4;
- a demonstration in Section 3 of the expressivity of these grammars by showing how to express the layout rules of ISWIM, Miranda, Haskell, and Python in terms of these grammars;
- a development in Section 5 of GLR and LR(k) parsing algorithms for these grammars, which is possible by a careful factoring of item sets into state and indentation sets; and
- a demonstration in Section 6 of the practical performance of these parsing techniques relative to existing *ad hoc* techniques.

Section 7 of this paper reviews related work. Section 8 concludes.

Note that Section 5 of this paper assumes a fair amount of familiarity with standard parsing techniques and in particular the work by Knuth [1965] on LR(k) parsing. Other than Section 5, however, this paper assumes only a basic knowledge of context-free grammars.

2. The Basic Idea

In order to support indentation-sensitive parsing, we use a modification of traditional, context-free grammars. We parse over a sequence of terminals where every terminal is annotated with the column at which it occurs in the source code. We call this its indentation. During parsing, we also annotate each non-terminal with an indentation. The grammar specifies a numerical relation that the indentation of each non-terminal must have with the indentation of its immediate children. These relations are usually chosen so the indentation of a non-terminal is the minimum column at which any token in the non-terminal is allowed to occur. Thus, the indentation of a non-terminal usually coincides with the intuitive notion of how far a block of code is indented. Formally, however, the indentation of a non-terminal has no meaning other than that it must appropriately relate to the indentation of the non-terminal's children.

We call these grammars indentation-sensitive context-free grammars (IS-CFG) to contrast them with traditional indentation-insensitive context-free grammars (II-CFG). Section 3 gives examples of IS-CFGs for real world languages, and Section 4 formally defines this class of grammars.

As a simple example, we may write $A \rightarrow ' (^=A^>) '=$ to mean that $($ and $)$ must be at the same indentation as the A on the left of the production arrow but the A on the right must be at a greater indentation. We may also write $A \rightarrow ' [^{\geq}A^>] '^{\geq}$ to mean the same except that $[$ and $]$ must be at a greater or equal indentation than the A on the left of the production arrow. In addition, we may write $A \rightarrow A^=A^=$ to mean that the indentation of both occurrences of A on the right of the production must be at indentations equal to that of the A on the left of the production. Combined with the production $A \rightarrow \varepsilon$, these form a grammar of nested parentheses and square brackets. In that grammar, matching parentheses must align vertically, and things enclosed in parentheses must be indented more than the parentheses are indented. Things enclosed in square brackets merely must be indented more than the surrounding code. Figure 1 shows examples of parse trees for this grammar on the words $(^1 [^4 (^5 ^5 ^7) ^1$ and $(^1 [^8 (^6 ^6 [^8] ^9] ^4 (^3 ^3) ^1$ where we write X^i to mean that i is the indentation of X . In these parse trees, take particular note of how the indentations of the non-terminals




Figure 1. Example IS-CFG parse trees for $(^1 [^4 (^5 ^5 ^7) ^1$ and $(^1 [^8 (^6 ^6 [^8] ^9] ^4 (^3 ^3) ^1$ respectively.

and terminals relate according to the indentation relations specified in the grammar.

In general, we write a production as $A \rightarrow X_1^{\triangleright_1} X_2^{\triangleright_2} \dots X_n^{\triangleright_n}$, where $\triangleright_1, \triangleright_2, \dots, \triangleright_n$ are relations between indentations, to mean that the indentation of A relates to the indentations of each X_1, X_2, \dots, X_n according to $\triangleright_1, \triangleright_2, \dots, \triangleright_n$. That is to say, this production requires $j_1 \triangleright_1 i, j_2 \triangleright_2 i, \dots, j_n \triangleright_n i$ if i is the indentation of A and j_1, j_2, \dots, j_n are the respective indentations of X_1, X_2, \dots, X_n .

While in principle any set of indentation relations can be used, we restrict ourselves to the relations $=$, $>$, \geq , and \otimes . The $=$, $>$, and \geq relations have their usual meanings. The \otimes relation is $\{(i, j) \mid i, j \in \mathbb{N}\}$ and effectively disassociates the indentation of a child from that of its parent.

Indentation-sensitive languages typically have forms that require the first token of a subexpression to be at the same indentation as the subexpression itself even though the non-terminal for that subexpression does not normally require this. Thus for every non-terminal or terminal, X , we introduce the non-terminal $|X|$ that is identical to X except that its indentation is always equal to the indentation of its first token. This is merely syntactic sugar, as we can introduce the production $|a| \rightarrow a^=$ for each terminal a and the production $|A| \rightarrow |X_m|^= X_{m+1}^{\triangleright_{m+1}} \dots X_n^{\triangleright_n}$ for each production $A \rightarrow X_1^{\triangleright_1} X_2^{\triangleright_2} \dots X_n^{\triangleright_n}$ where $n \geq 1$ and each $m \leq n$ such that X_1, X_2, \dots, X_{m-1} are all nullable. For example, with the above grammar, $|A|$ would have the productions $|A| \rightarrow ' (^=A^>) '=$ and $|A| \rightarrow ' [^{\geq}A^>] '^{\geq}$ from the first two productions for A . It would also have $|A| \rightarrow |A^=|^=$ and $|A| \rightarrow |A^=|^=$ from the third production for A . By replacing \triangleright_m with $=$, the first symbol of each production is forced to have the same indentation as the non-terminal on the left of the production. By transitivity, this is the indentation of the first token in the non-terminal. This is a straightforward, mechanical transformation requiring no input from the user.

3. Indentation-Sensitive Languages

Despite this system's simplicity, it can express a wide array of layout rules. This section demonstrates this by presenting the layout rules of several languages in terms of IS-CFGs. Where possible, we use the non-terminal names from the original grammar of each language. Though not shown here, sketches for other indentation-sensitive languages have been constructed for occam,¹ Orwell, Curry, Habit, and SRFI-49.

¹ The additional indentation relation $\{(i+2, i) \mid i \in \mathbb{N}\}$ is required by occam as it has forms that require increasing indentation by exactly 2.

3.1 ISWIM and Miranda

With ISWIM, Landin [1966] introduced several innovations still used in programming languages today. Among these is the use of indentation to indicate program structure. Landin defines an “offside rule”² that specifies that:

The southeast quadrant that just contains the phrase’s first symbol must contain the entire phrase, except possibly for [parenthesized] subsegments. [Landin 1966]

This rule requires that all the tokens in a particular non-terminal be at a column that is at least as far right (i.e., in the southeast quadrant) as the column of the first token.

The IS-CFG for ISWIM is similar to the II-CFG for ISWIM that ignores the offside rule except that we annotate each production of the IS-CFG with appropriate indentation relations. As an example, consider the productions in Figure 2 where we have applied the offside rule to only the right-hand side of *where* clauses.³ Note that the productions for $|expr|$ are created automatically from the productions for $expr$ as specified in Section 2. We do likewise for the terminals.⁴ Using this grammar to parse the expression

```
x + v where
x = -(
y + z) + w
```

results in the parse tree in Figure 3.

For productions that do not involve the offside rule, such as the productions for addition and negation, we annotate the non-terminals with $=$ and the terminals with \geq . This means that these productions do not change the current indentation and terminals are allowed at any column greater than or equal to the current indentation. For example, the expressions for $x + v$, $(y + z)$, and all the variable references are all at the same indentation as their respective parents. For those expressions, no special indentation rule is in effect and we simply use $=$ to propagate the indentation from parent to child. The terminals use \geq so they can be at that or any greater indentation. If a non-terminal is too far left (e.g., if w was at column 5), then the \geq constraint is violated and the code would be rejected. This is a common pattern that we will see in other grammars as well.

Next, consider forms that involve the offside rule. Wherever a non-terminal A should trigger the offside rule, we simply use $|A| \geq$ instead of $A =$. An example is the right-hand side of a *where* clause where we use $|expr| \geq$ instead of $expr =$. Because we use \geq instead of $=$, the right-hand side is allowed to be at a greater indentation than its parent. We see this in the example parse tree where the right-hand side of the *where* clause has an indentation of 6 while its parent has an indentation of 1. If the top-level $expr$ were at indentation 7, however, such a parse with a right-hand side at indentation 6 would be rejected. We also use $|expr|$ instead of $expr$ so that the first token of the right-hand side of the *where* clause has the same indentation as the entire right-hand side. As a consequence, in the example parse tree, the indentations on the path from the right-hand side of the *where* to the $'-'$ are all exactly 6.

Finally, since parenthesized expressions are exempt from the indentation constraints of the surrounding code, the production for parenthesized expressions uses $expr^\circ$ instead of $expr =$. This frees

Productions written by the user:

```
expr → expr = 'where' ≥ ID ≥ '=' ≥ |expr| ≥
expr → expr = '+' ≥ expr =
expr → '-' ≥ expr =
expr → '(' ≥ expr⊙ ')' ≥
expr → ID ≥
```

Productions added by desugaring:

```
|expr| → |expr| = 'where' ≥ ID ≥ '=' ≥ |expr| ≥
|expr| → |expr| = '+' ≥ expr =
|expr| → '-' ≥ expr =
|expr| → '(' ≥ expr⊙ ')' ≥
|expr| → ID =
```

```
'=' → '=' =      '+' → '+' =      '-' → '-' =
'(' → '(' =      ')' → ')' =      ID → ID =
'where' → 'where' =
```

Figure 2. IS-CFG productions for ISWIM.

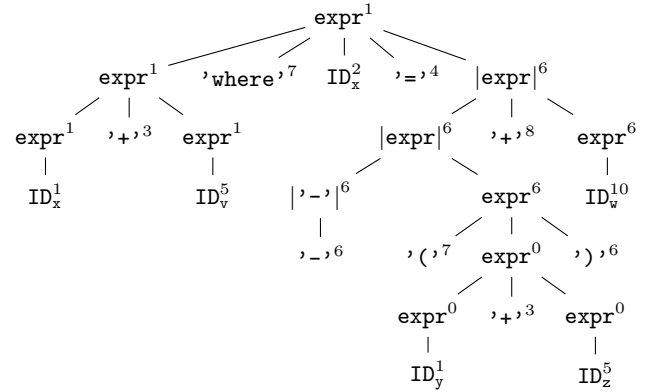


Figure 3. IS-CFG parse tree for ISWIM.

the indentation of the expression inside the parentheses from any indentation constraints coming from the context of the expression. This is used in the example parse tree where the $y + z$ expression has an indentation of 0 even though the $(y + z)$ expression has an indentation of 6. Uses of the offside rule inside the parentheses still have effect, of course.

The Miranda language uses the same offside rule as ISWIM except for two differences. The first is that expressions inside parentheses *are* subject to indentation constraints imposed by the context outside the parentheses. The production for parenthesized expressions thus uses $=$ instead of \circ and is simply:

```
expr → '(' ≥ expr = ')' ≥
```

The second difference is that the language specification adopts the notational convention that for any non-terminal x ,

$x(;$ means that x is followed by an optional semicolon and is subject to the offside rule..., so that every token of x must lie below or to the right of the first. Provided the layout makes it clear where x terminates, the trailing semicolon may be omitted. [Turner 1989, §25]

This notation is easily handled by introducing, for each non-terminal A , the non-terminal $A(;$ and the two productions $A(;) \rightarrow |A| \geq$ and $A(;) \rightarrow |A| \geq ';'.$

²Landin spells it as both “offside” and “off-side”. We adopt the spelling consistent with the rule from sports.

³Landin is not clear about whether his offside rule applies to all syntactic forms. His examples imply that it does not. For the sake of example, we also simplify *where* and omit the multiple bindings case. That case uses the same techniques as Haskell’s statement blocks.

⁴Of course, the productions generated for the terminals $'='$, $'+'$, $'-'$, $'('$ and $'where'$ are unreachable from the rest of the grammar and can be omitted.

```

L (<n>:ts) (m:ms) = ';' : (L ts (m:ms)) if m = n
                  = '}' : (L (<n>:ts) ms) if n < m
L (<n>:ts) ms = L ts ms
L ({n}:ts) (m:ms) = '{' : (L ts (n:m:ms)) if n > m
L ({n}:ts) [] = '{' : (L ts [n]) if n > 0
L ({n}:ts) ms = '{' : '}' : (L (<n>:ts) ms)
L ('}':ts) (0:ms) = '}' : (L ts ms)
L ('}':ts) ms = parse-error
L ('{':ts) ms = '{' : (L ts (0:ms))
L (t :ts) (m:ms) = '}' : (L (t:ts) ms)
                  if m ≠ 0 and parse-error(t)
L (t :ts) ms = t : (L ts ms)
L [] [] = []
L [] (m:ms) = '}' : L [] ms if m ≠ 0

```

Figure 4. Haskell’s L function [Marlow (ed.) 2010, §10.3].

3.2 Haskell

3.2.1 Language

Haskell uses a more sophisticated offside rule than does ISWIM. Indentation-sensitive blocks (e.g. the bodies of `do`, `case`, or `where` expressions) are made up of one or more statements or clauses that not only are indented relative to the surrounding code but also are indented to the same column as each other. Thus, lines that are more indented than the block continue the current clause, lines that are at the same indentation as the block start a new clause, and lines that are less indented than the block are not part of the block. In addition, semicolons (;) and curly braces ({ and }) can explicitly separate clauses and delimit blocks, respectively. Explicitly delimited blocks are exempt from indentation restrictions arising from the surrounding code.

While the indentation rules of Haskell are intuitive to use in practice, the way that they are formally expressed in the Haskell language specification [Marlow (ed.) 2010, §10.3] is not nearly so intuitive. The indentation rules are specified in terms of both the lexer and an extra pass between the lexer and the parser. Roughly speaking, the lexer inserts special `{n}` tokens where a new block might start and special `<n>` tokens where a new clause within a block might start. The extra pass then translates these tokens into explicit semicolons and curly braces.

The special tokens are inserted according to the following rules:

- If a `let`, `where`, `do`, or `of` keyword is not followed by the lexeme `{`, the token `{n}` is inserted after the keyword, where `n` is the indentation of the next lexeme if there is one, or 0 if the end of file has been reached.
- If the first lexeme of a module is not `{` or `module`, then it is preceded by `{n}` where `n` is the indentation of the lexeme.
- Where the start of a lexeme is preceded only by white space on the same line, this lexeme is preceded by `<n>`, where `n` is the indentation of the lexeme, provided that it is not, as a consequence of the first two rules, preceded by `{n}`. [Marlow (ed.) 2010, §10.3]

Between the lexer and the parser, an indentation resolution pass converts the lexeme stream into a stream that uses explicit semicolons and curly braces to delimit clauses and blocks. The stream of tokens from this pass is defined to be `L tokens []` where `tokens` is the stream of tokens from the lexer and `L` is the function in Figure 4. Thus the context-free grammar has to deal with only semicolons and curly braces. It does not deal with layout.

This `L` function is fairly intricate, but the key clauses are the ones dealing with `<n>` and `{n}`. After a `let`, `where`, `do`, or `of` keyword, the lexer inserts a `{n}` token. If `n` is a greater indentation than the current indentation, then the first clause for `{n}` executes,

```

case → 'case'> exp = 'of'> altBlock =

-- Explicitly delimited blocks
altBlock → '{'> alts@ '}',@

-- Layout-delimited blocks
altBlock → altLayout>
altLayout → altLayout = |alts| =
altLayout → |alts| =

-- Clause sequences
alts → alt =
alts → alts = ';'> alt =

```

Figure 5. Grammatical productions for `case`.

an open brace (`{`) is inserted, and the indentation `n` is pushed on the second argument to `L` (i.e., the stack of indentations). If a line starts at the same indentation as the top of the stack, then the first clause for `<n>` executes and a semicolon (;) is inserted to start a new clause. If it starts at a smaller indentation, then the second clause for `<n>` executes and a close brace (`}`) is inserted to close the block started by the inserted open brace. Finally, if the line is at a greater indentation, then the third clause executes, no extra token is inserted, and the line is a continuation of the current clause. The effect of all this is that `{`, `;`, and `}` tokens are inserted wherever layout indicates that blocks start, new clauses begin, or blocks end, respectively. The other clauses in `L` handle a variety of other edge cases and scenarios.

Note that `L` uses `parse-error` to signal a parse error, but uses `parse-error(t)` as an oracle that predicts the future behavior of the parser that runs after `L`. Specifically,

if the tokens generated so far by `L` together with the next token `t` represent an invalid prefix of the Haskell grammar, and the tokens generated so far by `L` followed by the token “`}`” represent a valid prefix of the Haskell grammar, then `parse-error(t)` is true. [Marlow (ed.) 2010, §10.3]

This handles code such as

```
let x = do f; g in x
```

where the block starting after the `do` needs to be terminated before the `in`. This requires knowledge about the parse structure in order to be handled properly, and thus `parse-error(t)` is used to query the parser for this information.

In addition to the operational nature of this definition, the use of the `parse-error(t)` predicate means that `L` cannot run as an independent pass; its execution must interact with the parser. In fact, the Haskell implementations GHC [GHC 2011] and Hugs [Jones 1994] do not use a separate pass for `L`. Instead, the lexer and parser share state consisting of a stack of indentations. The parser accounts for the behavior of `parse-error(t)` by making close braces optional in the grammar and appropriately adjusting the indentation stack when braces are omitted. The protocol relies on “some mildly complicated interactions between the lexer and parser” [Jones 1994] and is tricky to use. While preparing the parser in Section 6, we found that even minor changes to the error propagation of the parser affected whether syntactically correct programs were accepted by this style of parser.

While we may believe the correctness of these parsers based on their many years of use and testing, the significant and fundamental structural differences between their implementation and the language specification are troubling.

3.2.2 Grammar

Haskell's layout rule is more complicated than those of ISWIM and Miranda, but is also easily specified as an IS-CFG. By using an IS-CFG there is no need for an intermediate L function, and the lexer and parser can be cleanly separated into self-contained passes. The functionality of `parse-error(t)` is simply implicit in the structure of the grammar.

Figure 5 shows example productions for `case` expressions. For productions that do not change the indentation, we annotate non-terminals with a default indentation relation of `=` and terminals with a default indentation relation of `>`. We use `>` instead of `≥` because Haskell distinguishes tokens that are at an indentation equal to the current indentation from tokens that are at a strictly greater indentation. The former start a new clause while the latter continue the current clause.

In Haskell, a block can be delimited by either explicit curly braces or use of the layout rule. In Figure 5, this is reflected by the two different productions for `altBlock`. If `altBlock` expands to `'{'> alts* }*`, then the `*` relation allows `alts` to not respect the indentation constraints from the surrounding code.⁵ Since Haskell's layout rule allows closing braces to occur at any column, we use `*` instead of the usual `>` on `'}'`. On the other hand, if `altBlock` expands to `altLayout>`, then the `>` relation increases the indentation. In the productions for `altLayout`, the use of `|alts|` instead of `alts` ensures that the first tokens of the `alts` all align to the same column. Note that within an `alts`, each `alt` must be separated by a semicolon (`;`). Thus, because `altLayout` refers to `alts` instead of `alt`, each instance of `alt` can be separated using either layout or a semicolon. When using curly braces to explicitly delimit a block, semicolons must always be used.

Other grammatical forms that use the layout rule follow the same general pattern as `case` with only minor variation to account for differing base cases (e.g., `let` uses `decl` in place of `alt`) and structures (e.g., a `do` block is a sequence of `stmt` ending in an `exp`).

A subtlety of Haskell's layout rule is that tokens on the same line as, but after, a closing brace may not have to respect the current indentation. This is because the L function considers the indentation of only the first token of a line (i.e., where `<n>` is inserted) and tokens after a `let`, `where`, `do` or `of` keyword (i.e., where `{n}` is inserted). One might view this as an artifact of how the language specification uses L to define layout, but this aspect of Haskell's layout rule is still expressible by having the lexer annotate tokens whose indentation is to be ignored with an indentation of infinity.⁶ Since terminals have an indentation relation of `>`, the infinite indentation of these tokens will always match. We have the lexer handle this instead of the parser because it is the linear order of tokens instead of the grammatical structure of the syntax that controls what tokens are indentation sensitive. For example, the token after the `do` keyword is indentation sensitive regardless of the structure of the expression following the `do`. This requires the lexer to maintain a bit of extra state indicating whether we are at the start of a line or after a `let`, `where`, `do` or `of` keyword, but this is a fairly light requirement as the lexer is presumably already tracking state to determine the column of each token.

Finally, GHC also supports an alternative indentation rule that is enabled by the `RelaxedLayout` extension. It allows opening braces to be at any column regardless of the current indentation [GHC 2011, §1.5.2]. This is easily implemented by changing the first production for `altBlock` to be:

$$\text{altBlock} \rightarrow \text{'{' }^* \text{alts}^* \text{'}' }^*$$

⁵ This assumes no tokens are at column 0, which we reserve for this purpose.

⁶ Of course, column information for error reporting should still use the actual position of the token.

3.3 Python

3.3.1 Language

Python represents a different approach to specifying indentation sensitivity. It is explicitly line oriented and features `NEWLINE` in its grammar as a terminal that separates statements. The grammar uses `INDENT` and `DEDENT` tokens to delimit indentation-sensitive forms. An `INDENT` token is emitted by the lexer whenever the start of a line is at a strictly greater indentation than the previous line. Matching `DEDENT` tokens are emitted when a line starts at a lesser indentation.

In Python, indentation is used only to delimit statements, and there are no indentation-sensitive forms for expressions. This, combined with the simple layout rules, would seem to make parsing Python much simpler than for Haskell, but Python has line joining rules that complicate matters.

Normally, each new line of Python code starts a new statement. If, however, the preceding line ends in a backslash (`\`), then the current line is “joined” with the preceding line and is a continuation of the preceding line. In addition, tokens on this line are treated as if they had the same indentation as the backslash itself.

Python's *explicit* line joining rule is simple enough to implement directly in the lexer, but Python also has an *implicit* line joining rule. Specifically, expressions

```
in parentheses, square brackets or curly braces can be split
over more than one physical line without using backslashes.
... The indentation of the continuation lines is not important.
[Python, §2.1.6]
```

This means that `INDENT` and `DEDENT` tokens must not be emitted by the lexer between paired delimiters. For example, the second line of the following code should not emit an `INDENT` and the indentation of the third line should be compared to the indentation of the first line instead of the second line.

```
x = [
  y ]
z = 3
```

Thus, while the simplicity of Python's indentation rules is attractive, they contain hidden complexity that requires interleaving the execution of the lexer and parser.

3.3.2 Grammar

Though Python's specification presents its indentation rules quite differently from Haskell's specification, once we translate it to an IS-CFG, it shares many similarities with that of Haskell. The lexer still needs to produce `NEWLINE` tokens, but it does not produce `INDENT` or `DEDENT` tokens. As with Haskell, we start with a grammar where the non-terminals and terminals are annotated with indentation relations of `=` and `>`, respectively.

In Python, the only form that changes indentation is the `suite` non-terminal, which represents a block of statements contained inside a compound statement. For example, one of the productions for `while` is:

$$\text{while_stmt} \rightarrow \text{'while' } ^> \text{ test } ^{'} ^> \text{ suite } ^{=}$$

A `suite` has two forms. The first is for a single-line statement and is the same as with the standard Python grammar. The second is for multi-line statements. The following productions handle both of these two cases.

```
suite → stmt_list^ NEWLINE^
suite → NEWLINE^ block^
block → block^ |statement|^
block → |statement|^
```

When a `suite` is of the multi-line form (i.e., using the second production), the initial `NEWLINE` token ensures that the `suite` is on a separate line from the preceding header. The `block` inside a `suite` must then be at some indentation greater than the current indentation. Such a block is a sequence of `statement` forms that all start with their first token at the same column. In Python’s grammar, the productions for `statement` already include a terminating `NEWLINE`, so `NEWLINE` is not needed in the productions for `block`.

For implicit line joining, we employ the same trick as for parenthesized expressions in ISWIM and braces in Haskell. For any production that contains parentheses, square brackets or curly braces, we annotate the part contained in the delimiters with the \otimes indentation relation. Since the final delimiter is also allowed to appear at any column, we annotate it with \otimes . For example, one of the productions for list construction becomes:

`atom` \rightarrow ‘[’ \otimes `listmaker` \otimes ‘]’

There remain a few subtleties with Python’s line joining rules that we must address. First, as with Haskell, tokens after a closing delimiter can appear at any column. For example, the following code is validly indented according to Python’s rules:

```
while True:
    x = 1 + (
2) + 3
```

To handle this we use the same trick as for Haskell and annotate tokens that are not at the start of a line with an infinite indentation.

Second, while a lexer based on regular expressions can detect the start of a line and thus produce finite indentations for the first token of a line but infinite indentations for other tokens, it cannot detect matching parentheses to determine that `NEWLINE` tokens should be omitted inside delimited forms. Thus non-terminals that occur inside delimited forms need to allow the insertion of `NEWLINE` tokens at arbitrary locations. This may mean there have to be two forms of a non-terminal (i.e., for expressions inside versus outside a delimited form), but this is a fairly mechanical transformation that can be automated by the use of syntactic sugar similar to the syntactic sugar for $|A|$. Alternatively, it may be possible to use a grammar that does not use `NEWLINE` tokens at all and instead, like for Haskell, uses vertical alignment to delimit statements.

Finally, as with the standard Python parser, the lexer still handles the explicit line joining that is triggered by a line ending in a backslash (`\`). It gives the tokens of an explicitly joined line the same indentation as the backslash itself, and the backslash is not emitted as a token.

3.4 Conventions and syntactic sugar

In an IS-CFG, every symbol in every production must be annotated with an indentation relation. In many indentation-sensitive languages, however, productions often allow terminals to appear at any indentation greater than the current indentation but do not themselves change the current indentation. Thus we can simplify the job of writing an IS-CFG by adopting the convention that if a symbol on the right-hand side of a production is not explicitly annotated with an indentation relation, then it implicitly defaults to $=$ if it is a non-terminal and $>$ if it is a terminal. For example, with this convention, the only productions in Figure 5 that need explicit annotations are those for `altBlock`. All other productions simply use the defaults. Using this convention most productions in a grammar do not have to be annotated with indentation relations. They thus look like ordinary II-CFG productions, and only the forms that explicitly deal with indentation must be explicitly annotated.

In addition, just as II-CFGs often allow the use of alternation bars ($|$) or Kleene stars ($*$) to simplify writing grammars, it is often convenient to allow symbols on the right-hand side of a production

to be annotated with a composition of indentation relations. Thus we might write $A \rightarrow C^{>\otimes}$ instead of the more verbose

$A \rightarrow B^{\otimes}$
 $B \rightarrow C^{>}$

These conventions are merely notational conveniences and do not affect the fundamental theory.

4. Indentation-Sensitive Grammars

The formalism for IS-CFGs that this paper proposes is an extension of II-CFGs. Thus to review the standard definition of II-CFGs, recall that a grammar is a four-tuple $G = (N, \Sigma, \delta, S)$ where N is a finite set of non-terminal symbols, Σ is a finite set of terminal symbols, δ is a finite production relation, and $S \in N$ is the start symbol. The relation δ is a subset of $N \times (N \cup \Sigma)^*$, and we write $A \rightarrow X_1 X_2 \cdots X_n$ for a tuple $(A, X_1 X_2 \cdots X_n)$ that is an element of δ .

As a notational convention let A, B, C be elements of N , let a, b, c be elements of Σ , and let X, Y, Z be elements of $N \cup \Sigma$. Let U, V, W be elements of $(N \cup \Sigma)^*$, and u, v, w be elements of Σ^* .

We define a rewrite relation $(\Rightarrow) \subseteq (N \cup \Sigma)^* \times (N \cup \Sigma)^*$ such that $UAV \Rightarrow UX_1 X_2 \cdots X_n V$ iff $A \rightarrow X_1 X_2 \cdots X_n$. We define (\Rightarrow^*) as the reflexive, transitive closure of (\Rightarrow) .

The language recognized by a grammar is then defined as $\mathcal{L}(G) = \{w \in \Sigma^* \mid S \Rightarrow^* w\}$ and is the set of words reachable by the rewrite relation from the start symbol.

An IS-CFG is also a four-tuple, $G = (N, \Sigma, \delta, S)$, except that δ and S account for indentations. S is an element of $N \times \mathbb{N}$ and records the indentation of the initial non-terminal. The production relation, δ , is an element of $N \times ((N \cup \Sigma) \times \mathbb{I})^*$ where \mathbb{I} is the domain of indentation relations and each indentation relation is a subset of $\mathbb{N} \times \mathbb{N}$. In principle, these indentation relations can be any subset of $\mathbb{N} \times \mathbb{N}$, but for our purposes we restrict \mathbb{I} to the relations $=, >, \geq$ and \otimes .

Here and in the remainder of this paper, we restrict ourselves to finite indentations, but everything generalizes straightforwardly to languages with infinite indentations.

As a notational convention, let i, j and l be indentations and \triangleright be an indentation relation. For the sake of compactness, we adopt the notations X^i and X^{\triangleright} , respectively, for a pair of X and either an indentation i or an indentation relation \triangleright . Thus we write $A \rightarrow X_1^{\triangleright_1} X_2^{\triangleright_2} \cdots X_n^{\triangleright_n}$ for a tuple $(A, (X_1, \triangleright_1) (X_2, \triangleright_2) \cdots (X_n, \triangleright_n))$ that is an element of δ .

As with II-CFGs, we define a rewrite relation

$$(\Rightarrow) \subseteq ((N \cup \Sigma) \times \mathbb{N})^* \times ((N \cup \Sigma) \times \mathbb{N})^*$$

where $UAV \Rightarrow UX_1^{j_1} X_2^{j_2} \cdots X_n^{j_n} V$ iff $A \rightarrow X_1^{\triangleright_1} X_2^{\triangleright_2} \cdots X_n^{\triangleright_n}$ and $j_1 \triangleright_1 i, j_2 \triangleright_2 i, \dots, j_n \triangleright_n i$. The (\Rightarrow^*) relation, the language $\mathcal{L}(G)$, derivations, and parse trees are all defined as with II-CFGs except that they are in terms of this new rewrite relation.

Note that every II-CFG is encodable as an IS-CFG by translating every production $A \rightarrow X_1 X_2 \cdots X_n$ to $A \rightarrow X_1^{\otimes} X_2^{\otimes} \cdots X_n^{\otimes}$ and every word $a_1 a_2 \cdots a_m$ to $a_1^{i_1} a_2^{i_2} \cdots a_m^{i_m}$ with arbitrary $i_1, i_2, \dots, i_m \in \mathbb{N}$. Conversely, erasing the indentations and indentation relations in an IS-CFG results in an II-CFG. Note that translating from an II-CFG to an IS-CFG will not introduce ambiguities, but translating from an IS-CFG to an II-CFG might.

5. Parsing

Of course, a grammar is not practically useful if we cannot effectively parse with it. In this section, we show how to modify traditional parsing techniques for II-CFGs to handle IS-CFGs. We show this for both GLR and $LR(k)$ parsing. This can also be done for

CYK, SLR, LALR, GLL, and $LL(k)$, but we do not present those here as they are straightforward once the techniques for $LR(k)$ parsing are understood.

In order to derive GLR and $LR(k)$ parsing algorithms, we first prove a number of basic properties about indentation relations and IS-CFGs (Section 5.1). Then, we model IS-CFGs by using infinite II-CFGs (Section 5.2). Next, we consider traditional rewrite systems (Section 5.3) and parsing algorithms (Section 5.4) applied to this infinite II-CFG. Finally, we factor out the parts of these constructions representing indentations so that these algorithms can be expressed finitely (Sections 5.5, 5.6 and 5.7) and discuss some practical efficiency considerations (Section 5.8). The key insights here are, first, expressing the semantics of an IS-CFG in terms of an infinite II-CFG and, second, factoring the representation of item sets into a finite representation.

5.1 Basic properties

There are a few technical definitions and properties that we will use in our parsing algorithms. We present these without discussion.

Definition 1. A non-terminal A is *nullable* if $A^i \Rightarrow^* \epsilon$ for all i .

Lemma 2 (Composition of indentation relations). *Every finite sequence of compositions of elements from \mathbb{I} is one of $=$, \geq , \otimes , $\geq \otimes$, $>^n$, or $>^n \otimes$ for some $n \geq 1$.*

Proof. $=$ is a left and right identity under composition. \otimes is a left annihilator under composition. The compositions of $>$ with \geq and \geq with $>$ are both $>$. \square

Lemma 3 (Closure of indentation relations). *The closure of \mathbb{I} under finite sequences of composition and either finite or infinite sets of unions is the set of unions of one or more of $=$, \geq , \otimes , $\geq \otimes$, $>^n$, and $>^m \otimes$ for some $n, m \geq 1$ where each of these is in the union only once. We call this closure $\bar{\mathbb{I}}$ and as a notational convention let \triangleright be an element of $\bar{\mathbb{I}}$.*

Proof. By Lemma 2, every finite composition is of the form $=$, \geq , \otimes , $\geq \otimes$, $>^n$, or $>^m \otimes$. Since $>^n \triangleright >^{n'}$ if $n' \geq n$ and $>^m \otimes \triangleright >^{m'}$ if $m' \geq m$, at most one occurrence of $>^n$ and one occurrence of $>^m \otimes$ needs to be in the union. \square

Lemma 4 (Indentations of unique parses). *If $A^i \Rightarrow^* W$, then the set of all indentations i' such that $A^{i'} \Rightarrow^* W$ using the same sequence of productions is either the set \mathbb{N} , a singleton or the upper bounded set $\{i \mid i \leq n\}$ for some $n \in \mathbb{N}$. Furthermore, if it is a singleton or upper bounded set, then the maximum indentation is limited to be at most the maximum indentation in W .*

Proof. Consider the derivation as a parse tree. Each edge can be annotated with the indentation relation between parent and child nodes. The relation between the indentation of the root A^i and each leaf a^j is then a composition of the edges in the path from A^i to a^j . The possible values of i are the values compatible with every leaf indentation and the root's relation to them. By Lemma 2 every such relation is one of $=$, \geq , \otimes , $\geq \otimes$, $>^p$ or $>^q \otimes$ and for a particular leaf the compatible root indentations are thus either \mathbb{N} , a singleton, or $\{i \mid i \leq m\}$ for some $m \in \mathbb{N}$. The intersection of these over all the leaves in the parse tree is thus either the set \mathbb{N} , a singleton, or $\{i \mid i \leq n\}$ for some $n \in \mathbb{N}$. \square

Lemma 5 (Indentations of ambiguous parses). *If $A^i \Rightarrow^* W$, then the set of all indentations i' such that $A^{i'} \Rightarrow^* W$ using any sequence of productions is either the set \mathbb{N} or a finite subset of \mathbb{N} . Moreover, the finite subsets of \mathbb{N} are bounded by the maximum indentation in W .*

Proof. By Lemma 4 and the union over all parses. \square

5.2 Translating IS-CFGs to infinite II-CFGs

The first step of our approach is to model the IS-CFG by an infinite II-CFG. Of course we do not actually compute with this infinite grammar, but it provides a mathematical model from which we derive a computable parsing algorithm. Given an IS-CFG $G = (N, \Sigma, \delta, S)$, this II-CFG is $G' = (N', \Sigma', \delta', S')$ where:

$$\begin{aligned} N' &= N \times \mathbb{N} \\ \Sigma' &= \Sigma \times \mathbb{N} \\ \delta' &= \left\{ A^i \rightarrow X_1^{j_1} X_2^{j_2} \dots X_n^{j_n} \right. \\ &\quad \left| A \rightarrow X_1^{\triangleright_1} X_2^{\triangleright_2} \dots X_n^{\triangleright_n} \in \delta, \right. \\ &\quad \left. i, j_1, j_2, \dots, j_n \in \mathbb{N}, \right. \\ &\quad \left. j_1 \triangleright_1 i, j_2 \triangleright_2 i, \dots, j_n \triangleright_n i \right\} \\ S' &= S \end{aligned}$$

This grammar has an infinite number of non-terminals, terminals and productions per non-terminal, but we still limit derivations to finite lengths.

Note that traditional parsing algorithms on this grammar may not terminate due to the infinite size of G' , so we formally model G' as the limit of successive approximations where each approximation bounds the maximum indentation in any non-terminal, terminal or production to successively greater values. We gloss over this detail in the remainder of this paper.

Lemma 6 (Equivalence). $S \Rightarrow^* W$ for G iff $S' \Rightarrow^* W$ for G' .

Proof. By induction on the number of reductions and the fact that for all W and W' , $W \Rightarrow W'$ in G iff $W \Rightarrow W'$ in G' . \square

5.3 Rewrite system

In this subsection, we consider $LR(k)$ parsing in terms of a rewrite system that concisely specifies what it means for a parser to be $LR(k)$. In later subsections, we derive more conventional stack-based parsing algorithms. In this development, we closely follow the original presentation of $LR(k)$ parsing by Knuth [1965] with only minor changes to use current notational conventions.

Recall that an $LR(k)$ parser is one that always produces a rightmost derivation, and a rightmost derivation is one in which the rightmost non-terminal is always expanded before any other non-terminals. The symbols resulting from such an expansion step are called the handle. For example, if

$$S \Rightarrow^* UX^jV \Rightarrow UY_1^{l_1} \dots Y_p^{l_p}V \Rightarrow^* W$$

is a rightmost derivation, then a handle of $UY_1^{l_1} \dots Y_p^{l_p}V$ is $Y_1^{l_1} \dots Y_p^{l_p}$. Note that in order for this to be a rightmost derivation, V necessarily contains only elements of Σ , though U and $Y_1^{l_1} \dots Y_p^{l_p}$ may contain elements of both Σ and N .

Since an $LR(k)$ parser works from the result of a rightmost derivation back to the start symbol, $LR(k)$ parsing can be accomplished by iteratively searching for the handle of a string and performing the appropriate reduction. Again following Knuth [1965], given the infinite II-CFG $G' = (N', \Sigma', \delta', S')$ we construct the right-linear (and thus regular) grammar $G'' = (N'', N' \cup \Sigma', \delta'', S'')$ for recognizing prefixes that end in a handle. Here the non-terminals are

$$N'' = \left\{ \left[A^i, a_1^{j_1} a_2^{j_2} \dots a_k^{j_k} \right] \mid A^i \in N', a_1^{j_1}, a_2^{j_2}, \dots, a_k^{j_k} \in \Sigma' \right\}$$

and represent the part of the string that contains the handle. The $a_1^{j_1} a_2^{j_2} \dots a_k^{j_k}$ track the k terminals expected after the handle and

are the lookahead. For each $A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} X_{m+1}^{j_{m+1}} \dots X_n^{j_n}$ in δ' and each $u = a_1^{l_1} a_2^{l_2} \dots a_k^{l_k}$, we include the following in δ'' :

$$\begin{aligned} [A^i, u] &\rightarrow X_1^{j_1} \dots X_n^{j_n} u \\ [A^i, u] &\rightarrow X_1^{j_1} \dots X_m^{j_m} [X_{m+1}^{j_{m+1}}, v] \\ &\text{for each } v \in H_k(X_m^{j_m} \dots X_n^{j_n} u) \end{aligned}$$

Here, u and v are the lookaheads expected by the parser. $H_k(W)$ computes such lookaheads by computing the k -length prefixes of $\mathcal{L}(W)$ and is defined as

$$H_k(W) = \left\{ a_1^{l_1} a_2^{l_2} \dots a_k^{l_k} \mid W \Rightarrow^* a_1^{l_1} a_2^{l_2} \dots a_k^{l_k} U \right\}$$

where the reduction relation \Rightarrow^* is for G' .

The intuition here is that the first production expands to the handle along with a lookahead string and the second production expands to an intermediate non-terminal that in turn eventually expands to the handle.

Since this grammar is regular, it can be implemented by a state machine [Brzowski 1964], which leads to the following rewrite based algorithm for parsing.

Algorithm 7. Given input string W , if $S = W$ then stop and accept the string. Otherwise, find all prefixes of W that match the regular language $\mathcal{L}(G')$. If there are no such matches, then reject the string. Otherwise, non-deterministically choose one of the matches, and let the last production of the match be

$$[A^i, u] \rightarrow X_1^{j_1} X_2^{j_2} \dots X_n^{j_n} u$$

Replace this occurrence of $X_1^{j_1} X_2^{j_2} \dots X_n^{j_n}$ in W with A^i as it is a handle of W , and repeat this algorithm with the new value of W .

Note that this algorithm is non-deterministic and accepts the word if any path through the algorithm accepts the word. We allow this because productions in G (e.g., $A \rightarrow B^\circ$ and $A \rightarrow C^>$) may produce multiple productions when translated to G' (e.g., all $A^i \rightarrow B^j$ and $A^i \rightarrow C^k$ such that $k > i$). These may introduce ambiguities in G' even when there are no ambiguities in G . Once we convert to a finite version of the parsing algorithm, we will eliminate this non-determinism.

5.4 Parsing with stacks

Of course, rewriting the entire string and restarting the automaton from the start as done in Algorithm 7 is inefficient. Instead, we can save a trace of the states visited. When a handle is reduced, we rewind to the state just before the first symbol of the handle and proceed from there. This is the essential idea behind the traditional LR(k) parser development by Knuth [1965]. We apply this idea to our infinite II-CFG to obtain the following construction.

We begin with the notion of an item. We denote an item by

$$[A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j_{m+1}} \dots X_n^{j_n}; u]$$

where $A^i \rightarrow X_1^{j_1} \dots X_n^{j_n}$ is a production in G' and $u \in (\Sigma')^k$ is the lookahead. The algorithm maintains a stack of sets of items $S_0 S_1 \dots S_n$ where S_n is the top element of the stack. We use the notation $S_0 S_1 \dots S_n \mid a_1 a_2 \dots a_k w$ to denote that $S_0 S_1 \dots S_n$ is the current stack and $a_1 a_2 \dots a_k w$ is the input remaining to be consumed by the parser.

To parse a word w , we start with the configuration

$$S_0 \mid w \dashv_1^{i_0} \dashv_2^{i_0} \dots \dashv_k^{i_0}$$

where $S_0 = \left\{ [\hat{S} \rightarrow \bullet S'; \dashv_1^0 \dashv_2^0 \dots \dashv_k^0] \right\}$. We let \hat{S} be a fresh non-terminal and $\dashv_1, \dashv_2, \dots, \dashv_k$ be fresh terminals that pad the string

to have at least k tokens of lookahead. We then run the following parsing algorithm.

Algorithm 8. Given configuration $S_0 S_1 \dots S_n \mid a_1^{i_1} a_2^{i_2} \dots a_k^{i_k} w$, if $[\hat{S} \rightarrow S \bullet; u] \in S_n$ and $a_1^{i_1} a_2^{i_2} \dots a_k^{i_k} w = \dashv_1^{i_0} \dashv_2^{i_0} \dots \dashv_k^{i_0}$, then accept. Otherwise:

1. Compute the closure, S' , of S_n where S' is the least set of items satisfying the recurrence

$$\begin{aligned} S' = S_n \cup \left\{ [X_{m+1}^{j_{m+1}} \rightarrow \bullet Y_1^{l_1} \dots Y_p^{l_p}; v] \right. \\ \left. \mid [A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j_{m+1}} \dots X_n^{j_n}; u] \in S_n, \right. \\ \left. X_{m+1}^{j_{m+1}} \rightarrow Y_1^{l_1} \dots Y_p^{l_p} \in \delta' \right. \\ \left. v \in H_k(X_{m+1}^{j_{m+1}} \dots X_n^{j_n} u) \right\} \end{aligned}$$

2. Compute the acceptable lookahead set \mathcal{K} where

$$\begin{aligned} \mathcal{K} = \left\{ v \mid [A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j_{m+1}} \dots X_n^{j_n}; u] \in S', \right. \\ \left. v \in H_k(X_{m+1}^{j_{m+1}} \dots X_n^{j_n} u) \right\} \end{aligned}$$

3. For each production $A^i \rightarrow X_1^{j_1} \dots X_n^{j_n}$ in δ' , compute the acceptable lookahead set $\mathcal{K}(A^i \rightarrow X_1^{j_1} \dots X_n^{j_n})$ where

$$\begin{aligned} \mathcal{K}(A^i \rightarrow X_1^{j_1} \dots X_n^{j_n}) = \\ \left\{ u \mid [A^i \rightarrow X_1^{j_1} \dots X_n^{j_n} \bullet; u] \in S' \right\} \end{aligned}$$

4. Let $GOTO(S, Z^l) =$

$$\begin{aligned} \left\{ [A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} X_{m+1}^{j_{m+1}} \bullet X_{m+2}^{j_{m+2}} \dots X_n^{j_n}; v] \right. \\ \left. \mid [A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j_{m+1}} X_{m+2}^{j_{m+2}} \dots X_n^{j_n}; u] \in S, \right. \\ \left. X_{m+1}^{j_{m+1}} = Z^l \right\} \end{aligned}$$

and non-deterministically choose one of the following.

- (a) If $a_1^{i_1} a_2^{i_2} \dots a_k^{i_k} \in \mathcal{K}$, then do a shift action by looping back to the start of the algorithm with the new configuration

$$S_0 S_1 \dots S_n GOTO(S_n, a_1^{i_1}) \mid a_2^{i_2} \dots a_k^{i_k} w$$

- (b) If $a_1^{i_1} a_2^{i_2} \dots a_k^{i_k} \in \mathcal{K}(A^i \rightarrow X_1^{j_1} \dots X_n^{j_n})$ for some production $A^i \rightarrow X_1^{j_1} \dots X_n^{j_n}$, then do a reduce action by looping back to the start of the algorithm with the new configuration

$$S_0 S_1 \dots S_{n-m} GOTO(S_{n-m}, A^i) \mid a_1^{i_1} a_2^{i_2} \dots a_k^{i_k} w$$

5.5 Finite representations of stacks

Algorithm 8 contains both non-determinism and infinite sets. Here we depart from Knuth [1965] in order to eliminate these. Up to this point, our parser is simply a standard LR(k) parser, albeit on an infinite II-CFG, and we rely on the correctness of the standard LR(k) parsing algorithm for our correctness. From here forward, we ensure correctness by ensuring that our modified version of the algorithm models the same item sets as Algorithm 8, albeit using a more efficient representation.

As a first step, consider the sets of items that form the stack. Each item is of the form

$$[A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j_{m+1}} \dots X_n^{j_n}; u]$$

where $A^i \in N \times N$, $X_1^{j_1}, \dots, X_n^{j_n} \in (N \cup \Sigma) \times N$, and $u \in (\Sigma \times N)^k$. Observe that the indentations to the left of the bullet,

j_1, \dots, j_m , do not effect the parsing process, and multiple items that differ only in the values of j_1, \dots, j_m can therefore be represented by a single item. This reduces the state space slightly, but we can go further and also factor out j_{m+1}, \dots, j_n by observing that the algorithm preserves the following completeness property for item sets.

Definition 9 (Item-set completeness). We say that an item set, \mathcal{S} , is complete if for every $A \rightarrow X_1^{\triangleright_1} \dots X_n^{\triangleright_n} \in \delta$,

$$\left[A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j_{m+1}} \dots X_n^{j_n}; u \right] \in \mathcal{S}$$

implies that

$$\mathcal{S} \supseteq \left\{ \left[A^i \rightarrow X_1^{j_1} \dots X_m^{j_m} \bullet X_{m+1}^{j'_{m+1}} \dots X_n^{j'_n}; u \right] \mid j'_{m+1}, \dots, j'_n \in \mathbb{N}, j'_{m+1} \triangleright_{m+1} i, \dots, j'_n \triangleright_n i \right\}$$

This property is preserved by each loop through Algorithm 8, as stated in the following lemma.

Lemma 10. *If every item set on the stack is complete at the start of a loop through Algorithm 8, then every item set on the new stack at the start of the next loop is also complete.*

Proof. For every production $X_{m+1} \rightarrow Y_1^{\triangleright_1} \dots Y_p^{\triangleright_p}$, if the closure step adds the item $\left[X_{m+1}^{j_{m+1}} \rightarrow \bullet Y_1^{l_1} \dots Y_p^{l_p}; v \right]$, then by construction it adds all items of the form $\left[X_{m+1}^{j_{m+1}} \rightarrow \bullet Y_1^{l'_1} \dots Y_p^{l'_p}; v \right]$ where $l'_1 \triangleright_1 j_{m+1}, \dots, l'_p \triangleright_p j_{m+1}$. Thus this step of the algorithm preserves completeness.

The *GOTO* operation filters the set of items by requiring that $X_{m+1}^{j_{m+1}} = Z^l$, but afterwards $X_{m+1}^{j_{m+1}}$ is moved to the left of the bullet and is therefore no longer relevant to the completeness of the item set. Since j_{m+1} is related to j_{m+2}, \dots, j_n only indirectly through i , the remaining j_{m+2}, \dots, j_n are complete given the i that remain in the item set. \square

Since the item sets in the stack are all complete, we no longer need to represent the individual j_1, \dots, j_n . The only indentations we need to record are i , the indentation of the non-terminal on the left-hand side of the production, and the indentations in u , the expected lookahead. Thus, we can represent items that differ only in the values of j_1, \dots, j_n with

$$\left[A \rightarrow X_1^{\triangleright_1} \dots X_m^{\triangleright_m} \bullet X_{m+1}^{\triangleright_{m+1}} \dots X_n^{\triangleright_n}; L \right]$$

where $A \rightarrow X_1^{\triangleright_1} \dots X_n^{\triangleright_n} \in \delta$ and $L \subseteq \mathbb{N} \times (\Sigma \times \mathbb{N})^k$. We call L the indentation-and-lookahead set. Each element of L is a pair of i , the indentation for A , and u , the lookahead word.

With this factoring we can implement a GLR parser [Tomita 1985] by letting $k = 0$ and translating Algorithm 8 to use this representation of item sets. Since $k = 0$, the lookahead is always the empty string and L is simply a set of indentations. By examining the algorithm we can further determine that this set will always be either the set \mathbb{N} , a finite subset of \mathbb{N} , or the union of a finite subset of \mathbb{N} with the set of all elements of \mathbb{N} greater than j for some $j \in \mathbb{N}$. These sets are all finitely representable and thus computable. Finally, the standard technique of elaborating the possible item sets and the transitions between them can be used to construct a state machine for a push-down automaton that recognizes G . During this elaboration, we keep the set of indentations, L , abstract. This ensures that there are only finitely many item sets to be elaborated. At runtime, both the current state and the stack elements then simply store a reference to one of these pre-elaborated item sets along with a concrete indentation set. Any non-determinism remaining in the algorithm is handled using the standard stack representation used by GLR parsers.

5.6 Parsing LR(k) grammars

After having factored the representation of the item sets this far, we eliminate the non-determinism in the algorithm by restricting the algorithm to only LR(k) grammars. Recall that an LR(k) grammar is defined to have a unique rightmost derivation for any given word. Thus there is always a unique handle at each parsing step, and thus in the non-deterministic choice at the end of Algorithm 8 there is never more than one allowed choice. Otherwise, we have a shift/reduce or a reduce/reduce conflict, and the grammar is not an LR(k) grammar.

With the original representation of item sets, we did not apply this criterion because there might be multiple reductions that differ only in their indentations. For example, if we have $A \rightarrow B^{\otimes}$ in G , then both $A^1 \rightarrow B^1$ and $A^2 \rightarrow B^1$ are in G' . These could lead to spurious reduce/reduce conflicts. However, now that items are identified in terms of productions from δ (e.g., $A \rightarrow B^{\otimes}$) instead of productions from δ' (e.g., $A^1 \rightarrow B^1$), these conflicts no longer occur. Any remaining shift/reduce or reduce/reduce conflicts reflect a conflict in the original grammar, G , and are not artifacts of the translation to G' .

The last remaining source of ambiguity is the indentation-and-lookahead set, L , which can grow infinitely. To resolve this we change the representation of L from being a subset of $\mathbb{N} \times (\Sigma \times \mathbb{N})^k$ to being a subset of \mathbb{N} along with an element of $(\bar{\mathbb{N}} \times \bar{\mathbb{N}} \times \Sigma)^k \times \bar{\mathbb{N}}$.

We write such an item as

$$\left[A \rightarrow X_1^{\triangleright_1} \dots X_m^{\triangleright_m} \bullet X_{m+1}^{\triangleright_{m+1}} \dots X_n^{\triangleright_n}; I; (\bar{\triangleright}_{1,1}, \bar{\triangleright}_{1,2}, a_1) \dots (\bar{\triangleright}_{k,1}, \bar{\triangleright}_{k,2}, a_k) \bar{\triangleright}_{k+1} \right]$$

This represents the item

$$\left[A \rightarrow X_1^{\triangleright_1} \dots X_m^{\triangleright_m} \bullet X_{m+1}^{\triangleright_{m+1}} \dots X_n^{\triangleright_n}; L \right]$$

where $S = B^{i_0}$ and

$$L = \left\{ \left(i, a_1^{j_1} \dots a_k^{j_k} \right) \mid i \in I, \right. \\ \left. i \bar{\triangleright}_{1,1} l_1, l_1 \bar{\triangleright}_{2,1} l_2, \dots, l_{k-1} \bar{\triangleright}_{k,1} l_k, \right. \\ \left. j_1 \bar{\triangleright}_{1,2} l_1, j_2 \bar{\triangleright}_{2,2} l_2, \dots, j_k \bar{\triangleright}_{k,2} l_k, \right. \\ \left. l_k \bar{\triangleright}_{k+1} i_0 \right\}$$

The intuition behind this representation is best understood in terms of the sort of parse tree that could lead to the item that is trying to parse an A with lookaheads of $a_1 \dots a_k$. This situation is depicted in Figure 6. Note that there are no terminals between any of A , or a_1, a_2, \dots, a_k , and in the general case, we consider a node to be an ancestor of itself so some of $C_1^{l_1}, C_2^{l_2}, \dots, C_k^{l_k}$ or B^{i_0} may actually be the same node.

The lookahead token $a_1^{j_1}$ is in the lookahead only because A shares with it the ancestor C_1 at indentation l_1 . The $\bar{\triangleright}_{1,1}$ and $\bar{\triangleright}_{1,2}$ relations record the possible indentation relations between C_1 and, respectively, A and $a_1^{j_1}$. The second lookahead token $a_2^{j_2}$ is in the lookahead only because C_1 , the common ancestor of the item and the first lookahead token, also shares with a_2 the common ancestor C_2 at indentation l_2 . The $\bar{\triangleright}_{2,1}$ and $\bar{\triangleright}_{2,2}$ relations record the possible indentation relations between C_2 and, respectively, C_1 and a_2 . And so on, until we reach the final ancestor, C_k , at indentation l_k . This ancestor has a minimum indentation at which it can occur so we use $\bar{\triangleright}_{k+1}$ to record the indentation relation between C_k and B , the start terminal.

The lookahead computation, H_k , that is defined in Section 5.3 must of course be modified to account for this representation. We omit this because, while conceptually simple, its formal definition is fairly intricate.




Figure 6. The structure of lookahead tokens in a parse tree.

With this representation, the lookahead set is finitely represented so the only potentially infinite part remaining is the indentation set. However, as before, we can show that these sets are always either the set \mathbb{N} , a finite subset of \mathbb{N} , or the union of a finite subset of \mathbb{N} with the set of all elements of \mathbb{N} greater than j for some $j \in \mathbb{N}$. Indentation sets are thus finitely representable, and we have completely reduced the parsing algorithm to a finite representation.

We can now construct an $LR(k)$ parser just as we constructed the GLR parser at the end of Section 5.5. As before, to do this we translate Algorithm 8 to use the new item set representation, keep the indentation sets abstract while we elaborate the possible item sets to form the states of the push-down automaton, and at runtime augment the automaton states and stack entries with concrete indentation sets.

5.7 Correctness of item set representation

We must consider carefully the correctness of the parsing algorithm based on the item-set representation given in Section 5.6. At first glance, the representation looks like it could lead to lookaheads matching when they should not. Indeed, the following grammar of variable references (i.e., ID) and do blocks with vertically aligned statements shows how this can arise:

```

expr  → ID>
expr  → 'do'> |stmts|>
stmts → |expr|=
stmts → stmts= |expr|=

```

For example, consider the parse of the word 'do'¹ 'do'⁴ ID⁷ ID² which may come from the code:

```

do do x
  y

```

Note that y does not align with either the second `do` or the `x`, and thus this code should be rejected. Put another way, the valid lookaheads when looking ahead at the ID for y but before reducing x to `expr` are ID⁴, ID⁷, 'do'⁴, and 'do'⁷. The string should thus be rejected since ID², the token for y , is not in that set. But the $LR(1)$ lookahead set using the new representation is $\{((\geq, =, ID) >), ((\geq, =, 'do') >)\}$. Since the current indentation is 7 and there exist l such that $7 \geq l$, $2 = l$, and $l > 0$, the string will not be immediately rejected.

This representation thus appears to over approximate the set of indentations for a particular lookahead. This means that in the non-deterministic choice at the end of Algorithm 8 there could be more reductions possible than there should be. However, as we restrict ourselves to $LR(k)$ grammars, this turns out to not be a problem.

This is because there are only four cases where these spurious reductions occur:

Case 1. There should be *no* reductions or shifts possible, but the approximation makes *one* extra reduction possible.

In this case, the parser should reject the program, but will instead reduce and continue parsing. However, this case happens only when some other item set further up the stack also checks the lookahead tokens. Though the string is not rejected immediately, it will be rejected once we reach that point in the stack. In our example, once `x` reduces to `expr` and then `do x` reduces to an `expr` that finally reduces to `stmts`, the lookahead set will be $\{((=, =, ID) >), ((=, =, 'do') >)\}$. Since the indentation at that point is 4 but the next token is ID², the parser will reject the program. The program might not be rejected as soon as we expect, but it is eventually rejected. This case may arise even when the grammar is $LR(k)$.

Case 2. There should be *no* shifts or reductions possible, but the approximation makes *two or more* extra reductions possible.

The representation for lookaheads in Section 5.6 is designed so that every lookahead word that it represents comes from a valid parse. Thus if this representation generates multiple possible reductions, then there is some string that would also generate those multiple possible reductions with the original representation in Algorithm 8. In that case, the grammar is not $LR(k)$, and the grammar should be rejected by the parser generator.

Case 3. There should be *one* reduction or shift possible, but the approximation makes *one or more* extra reductions possible.

The same reasoning applies as in the preceding case.

Case 4. There should be *two or more* reductions or shifts possible.

Then the original grammar is not $LR(k)$, and the grammar should be rejected by the parser generator.

This means that the representation in Section 5.6 is valid for any grammar that is $LR(k)$, and furthermore we can detect when a grammar is not $LR(k)$ by using this representation.

5.8 An efficiency consideration

Note that each item in an item set may have a different set of indentations. For example, we may have an item set containing both of the following two items:

$$[A \rightarrow \bullet a^> b^=; I; u]$$

$$[A \rightarrow \bullet a^= b^=; I; u]$$

Even if they start with the same indentation set, I , after reading an a , the indentation sets for these items will be different from each other. For the first item, the indentation set will be restricted to indentations strictly greater than the indentation for a . For the second item, the indentations will be restricted to indentations equal to that of a . Thus, different items can have different indentation sets, and a naive factoring (e.g., sharing I between items in an item set) is insufficient. This does not preclude the possibility of a clever refactoring like the one done with the lookahead sets, but we have been unable to find such a factoring that works in all cases. Nevertheless, as a practical matter the following techniques seem to work well.

Observe that we need to keep only the indentation sets for items before the closure is taken. Items generated by the closure operation can be annotated with the $\bar{\triangleright}$ that can lead to them and this value can be incorporated into the lookahead check. In addition, when we can determine that some set of items will always have the same indentation set, we can represent them using a common




Figure 7. Benchmark results.

indentation set. These techniques reduce the number of indentation sets that must be passed from one state to another, and based on the experience implementing the Haskell parser in Section 6, the resulting number of such sets is usually one.

6. Implementation

In order to verify the real-world practicality of this parsing technique, we modified the Happy parser generator [Marlow and Gill 2009] to support the parsing techniques presented in this paper. The parser is LALR instead of $LR(k)$, but the techniques shown in Section 5 generalize straightforwardly to LALR. Note that this implementation is only a prototype and is only intended for testing the practical feasibility of parsing with IS-CFGs. In particular, no significant effort was put into optimizing the performance of the generated parser.

The Haskell parser from the `haskell-src` package [Marlow et al. 2011] uses techniques for implementing layout similar to those used by GHC. However, it is packaged as a standalone parser, and this makes it easy to isolate for benchmarking purposes. This parser was modified to use an IS-CFG instead of using shared state between the lexer and parser. Both the modified and unmodified versions were then run on the preprocessed source files from the `base` package [bas 2012]. Two modules (`GHC.Float` and `GHC.Constants`) could not be preprocessed due to missing header files. Of the 178 remaining Haskell modules, 85 cannot be parsed by the unmodified `haskell-src` parser due to syntactic extensions, such as rank-2 types, that are not supported by `haskell-src`. This left 93 source files that are parsable by the unmodified parser. All of these files parsed and produced the same parse tree when parsed using the IS-CFG based parser.

Figure 7 shows the parsing times of the modified parser relative to the unmodified parser when run on these files. These benchmarks were compiled using GHC version 7.0.3 with the `-O2`, `-opt1-static` and `-opt1-pthread` flags and were run on a 64-bit, 3.2GHz Xeon running Ubuntu 12.04 with 4GB of RAM. Timing measurements were collected using Criterion version 0.6.0.1 with 1000 samples per benchmark.

Note that the `haskell-src` parser is designed to run both the lexer and the parser simultaneously so they can share information about indentations. It is difficult to separate the execution of the parser from the execution of the lexer, so the lexing time is included in the times for both that parser and the IS-CFG based one.

As expected, the IS-CFG based parser runs in approximately linear time. It is geometrically on average 1.73 times slower than

the unmodified `haskell-src` parser. There is a slight upward trend in the factor by which the IS-CFG based parser is slower than the unmodified parser. This is primarily due to the fact that we are graphing the ratio between the performance of the modified and unmodified parsers and the unmodified parser has low-order performance overheads that are more significant on small inputs. Given that this is a prototype implementation with little optimization, the fact that the IS-CFG version is only one to three times slower than the standard `haskell-src` parser is promising. This overhead is likely due to the manipulation of the indentation sets as the representation of indentation sets is naive. Since in practice only certain sorts of sets are common (e.g., singletons and the set \mathbb{N}), an improved version could optimize for these sorts of sets. In addition, we could take advantage of the tokens with an indentation of infinity by adding a fast path through the parser that short circuits the indentation computations.

7. Related Work

The `uulib` parser library [Swierstra 2011] and the `indents` [Anklesaria 2012] and `indentparser` [Kurur 2012] extensions to the Parsec [Leijen and Martini 2012] parser library provide support for indentation-sensitive parsing. To the best of our knowledge there is no published, formal theory for the sort of indentation that these parsers implement. They are all combinator-based, top-down parsers and use some variation of threading state through a parser monad to track the current indentation.

Hutton [1992] describes an approach to parsing indentation-sensitive languages that is based on filtering the token stream. This idea is further developed by Hutton and Meijer [1996]. In both cases, the layout combinator searches the token stream for appropriately indented tokens and passes only those tokens to the combinator for the expression to which the layout rule applies. As each use of layout scans the remaining tokens in the input, this can lead to quadratic running time. Given that the layout combinator filters tokens before parsing occurs, this technique also cannot support subexpressions, such as parenthesized expressions in Python, that are exempt from layout constraints. Thus, this approach is incapable of expressing many real-world languages including ISWIM, Habit, Haskell, and Python.

Erdweg et al. [2012] propose a method of parsing indentation-sensitive languages by effectively filtering the parse trees generated by a GLR parser. The GLR parser generates all possible parse trees irrespective of layout. Indentation constraints on each parse node then remove the trees that violate the layout rules. For performance reasons, this filtering is interleaved with the execution of the GLR parser when possible. Aside from the fact that they require a GLR parser and thus generate parse trees that might not be used, a critical difference between their system and the one presented in this paper is that their indentation constraints are in terms of the set of tokens under a non-terminal whereas the system in this paper uses constraints between non-terminals and their immediate children. Thus, the two approaches look at the problem from different perspectives. Erdweg et al. [2012] do not consider the question of an $LR(k)$ parser.

Brunauer and Mühlbacher [2006] take a unique approach to specifying the indentation-sensitive aspects of a language. They use a scannerless grammar that uses individual characters as tokens and has non-terminals that take an integer counter as parameter. This integer is threaded through the grammar and eventually specifies the number of spaces that must occur within certain productions. The grammar encodes the indentation rules of the language by carefully arranging how this parameter is threaded through the grammar and thus how many whitespace characters should occur at each point in the grammar.

While encoding indentation sensitivity this way is formally precise, it comes at a cost. The YAML specification [Ben-Kiki et al. 2009] uses the approach proposed by Brunauer and Mühlbacher [2006] and as a result has about a dozen and a half different non-terminals for various sorts of whitespace and comments. With this encoding, the grammar cannot use a separate tokenizer and must be scannerless, each possible occurrence of whitespace must be explicit in the grammar, and the grammar must carefully track which non-terminals produce or expect what sorts of whitespace. The authors of the YAML grammar establish naming conventions for non-terminals that help manage this, but the result is still a grammar that is difficult to comprehend and even more difficult to modify.

While this approach bears some similarity to the technique proposed in this paper, a key difference is that their method uses the parameters of non-terminals to generate explicit whitespace characters and thus incurs a significant accounting overhead in the design of the grammar. On the other hand, the system presented in this paper operates at a higher level, using the parameter to indicate the column or indentation at which non-terminals and terminals should occur. This is a subtle distinction, but it has a profound impact. As shown in Section 3, layout rules are comparatively simple to encode this way, and as shown in Section 5, this formalism is amenable to traditional parsing techniques such as $LR(k)$ parsing.

Note that none of the systems reviewed above present an $LR(k)$ parsing algorithm. They use either top-down parsers or, in the case of Erdweg et al. [2012], a GLR parser.

8. Conclusion

This paper presents a grammatical formalism for indentation-sensitive languages. It is both expressive and easy to use. We derive provably correct GLR and $LR(k)$ parsers for this formalism. Though not shown here, CYK, SLR, LALR, GLL and $LL(k)$ parsers can also be constructed by appropriately using the key technique of factoring item sets. Experiments on a Haskell parser using this formalism show that the parser runs between one and three times slower than a parser using traditional *ad hoc* techniques for handling indentation sensitivity. Improvements in the handling of indentation sets may reduce this overhead. Using these techniques, the layout rules of a wide variety of languages can be expressed easily and parsed effectively.

Acknowledgments

Feedback from Andrew Tolmach, R. Kent Dybvig, Tim Sheard, Mark P. Jones, Nathan Collins, Steffen Lösch and the anonymous referees helped improve the presentation of this paper.

References

- base version 4.5.1.0, June 2012. URL <http://hackage.haskell.org/package/base/>.
- Sam Anklesaria. indents version 0.3.3, May 2012. URL <http://hackage.haskell.org/package/indents/>.
- Oren Ben-Kiki, Clark Evans, and Ingy döt Net. *YAML Ain't Markup Language (YAML) Version 1.2*, 3rd edition, October 2009. URL <http://www.yaml.org/spec/1.2/spec.html>.
- Leonhard Brunauer and Bernhard Mühlbacher. Indentation sensitive languages. Unpublished manuscript, July 2006. URL <http://www.cs.uni-salzburg.at/~ck/wiki/uploads/TCS-Summer-2006.IndentationSensitiveLanguages/>.
- Janusz A. Brzozowski. Derivatives of regular expressions. *Journal of the ACM (JACM)*, 11(4):481–494, October 1964. ISSN 0004-5411. doi: 10.1145/321239.321249.
- Sebastian Erdweg, Tillmann Rendel, Christian Kästner, and Klaus Ostermann. Layout-sensitive generalized parsing. In *Software Language Engineering*, Lecture Notes in Computer Science. Springer Berlin / Heidelberg, 2012. URL <http://sugarj.org/layout-parsing.pdf>. To appear.
- The Glorious Glasgow Haskell Compilation System User's Guide, Version 7.2.1*. The GHC Team, August 2011. URL http://www.haskell.org/ghc/docs/7.2.1/html/users_guide/.
- David Goodger. *reStructuredText Markup Specification*, January 2012. URL <http://docutils.sourceforge.net/docs/ref/rst/restructuredtext.html>. Revision 7302.
- John Gruber. *Markdown: Syntax*. URL <http://daringfireball.net/projects/markdown/syntax>. Retrieved on June 24, 2012.
- Michael Hanus (ed.). Curry: An integrated functional logic language (version 0.8.2). Technical report, March 2006. URL <http://www.informatik.uni-kiel.de/~curry/report.html>.
- HASP Project. The Habit programming language: The revised preliminary report, November 2010. URL <http://hasp.cs.pdx.edu/habit-report-Nov2010.pdf>.
- Graham Hutton. Higher-order functions for parsing. *Journal of Functional Programming*, 2(03):323–343, July 1992. doi: 10.1017/S0956796800000411.
- Graham Hutton and Erik Meijer. Monadic parser combinators. Technical Report NOTTCS-TR-96-4, Department of Computer Science, University of Nottingham, 1996.
- INMOS Limited. *occam programming manual*. Prentice-Hall international series in computer science. Prentice-Hall International, 1984. ISBN 978-0-13-629296-8.
- Mark P. Jones. The implementation of the Gofer functional programming system. Research Report YALEU/DCS/RR-1030, Yale University, New Haven, Connecticut, USA, May 1994.
- Donald E. Knuth. On the translation of languages from left to right. *Information and Control*, 8(6):607–639, December 1965. ISSN 0019-9958. doi: 10.1016/S0019-9958(65)90426-2.
- Piyush P. Kurur. indentparser version 0.1, January 2012. URL <http://hackage.haskell.org/package/indentparser/>.
- P. J. Landin. The next 700 programming languages. *Communications of the ACM*, 9(3):157–166, March 1966. ISSN 0001-0782. doi: 10.1145/365230.365257.
- Daan Leijen and Paolo Martini. parsec version 3.1.3, June 2012. URL <http://hackage.haskell.org/package/parsec/>.
- Simon Marlow and Andy Gill. *Happy User Guide*, 2009. URL <http://www.haskell.org/happy/doc/html/>. For Happy version 1.18.
- Simon Marlow, Sven Panne, and Noel Winstanley. *haskell-src* version 1.0.1.5, November 2011. URL <http://hackage.haskell.org/package/haskell-src>.
- Simon Marlow (ed.). *Haskell 2010 Language Report*, April 2010. URL <http://www.haskell.org/onlinereport/haskell2010/>.
- Egil Möller. *SRFI-49: Indentation-sensitive syntax*, May 2005. URL <http://srfi.schemers.org/srfi-49/srfi-49.html>.
- Python. *The Python Language Reference*. URL <http://docs.python.org/reference/>. Retrieved on June 26, 2012.
- S. Doaitse Swierstra. uulib version 0.9.14, August 2011. URL <http://hackage.haskell.org/package/uulib/>.
- Don Syme et al. *The F# 2.0 Language Specification*. Microsoft Corporation, April 2010. URL <https://research.microsoft.com/en-us/um/cambridge/projects/fsharp/manual/spec.html>. Updated April 2012.
- Masaru Tomita. *Efficient Parsing for Natural Language: A Fast Algorithm for Practical Systems*. Kluwer International Series in Engineering and Computer Science. Kluwer Academic Publishers, 1985. ISBN 978-0-89838-202-0.
- D. A. Turner. *Miranda System Manual*. Research Software Limited, 1989. URL <http://www.cs.kent.ac.uk/people/staff/dat/miranda/manual/>.
- Philip Wadler. An introduction to Orwell. Technical report, Programming Research Group at Oxford University, 1985.